

```
import pandas as pd
import numpy as np
import seaborn as sns
import matplotlib.pyplot as plt

!apt-get install openjdk-8-jdk -qq> /dev/null

# Download Apache Spark
!wget https://d1cdn.apache.org/spark/spark-3.4.1/spark-3.4.1-bin-hadoop3.tgz

# Extract the downloaded archive
!tar xf spark-3.4.1-bin-hadoop3.tgz

# specify where the files are
import os
os.environ["JAVA_HOME"] = "/usr/lib/jvm/java-8-openjdk-amd64"
os.environ["SPARK_HOME"] = "/content/spark-3.4.1-bin-hadoop3"

!pip install -q findspark
import findspark

findspark.init()

from pyspark.sql.functions import col

import pyspark
from pyspark import SparkContext

# Create a local SparkContext
sc = SparkContext('local', 'MyApp')

from pyspark.sql import SparkSession
spark = SparkSession.builder.master("local(*)").getOrCreate()

#spark.conf.set("spark.sql.repl.eagerEval.enabled", True)

spark

--2023-08-26 12:23:09-- https://d1cdn.apache.org/spark/spark-3.4.1/spark-3.4.1-bin-hadoop3.tgz
Resolving d1cdn.apache.org (d1cdn.apache.org)... 151.101.2.132, 2a04:4e42::644
Connecting to d1cdn.apache.org (d1cdn.apache.org)|151.101.2.132|:443... connected.
HTTP request sent, awaiting response... 200 OK
Length: 388341449 (370M) [application/x-gzip]
Saving to: 'spark-3.4.1-bin-hadoop3.tgz'

spark-3.4.1-bin-had 100%[=====] 370.35M 231MB/s in 1.6s

2023-08-26 12:23:11 (231 MB/s) - 'spark-3.4.1-bin-hadoop3.tgz' saved [388341449/388341449]

SparkSession - in-memory
SparkContext
Spark UI

Version
v3.4.1
Master
local
AppName
MyApp

from google.colab import drive
drive.mount('/content/drive')

Mounted at /content/drive

raw_df = spark.read.csv('/content/drive/MyDrive/Data Analytics/housdata.csv', header=True, inferSchema=True)
raw_df.show()
```

	CONTROL	AGE1	BEDRMS	PER	REGION	METRO3	LMED	FMR	L30	L50	L80	IPOV	BUILT	STATUS	TYPE	VCHRMV	VALUE	VACANCY	TENURE	NU
'100003130103'	87	3	2	'1'	'3'	71779	1095	17448	29071	45266	12956	2006	'1'	1	'-6'	50000	-6	'1'		
'100003130203'	70	3	1	'1'	'3'	71779	1095	15272	25446	39602	10292	2006	'1'	1	'-6'	238000	-6	'1'		
'100006370140'	48	4	4	'3'	'5'	53872	965	16555	27594	44150	22159	1985	'1'	1	'-6'	200000	-6	'1'		
'100006520140'	62	3	2	'3'	'5'	53872	861	13245	22076	35319	14370	1985	'1'	1	'-6'	175000	-6	'1'		
'100007130148'	30	2	2	'3'	'1'	61059	685	14662	24438	39103	14774	1980	'1'	1	'-6'	-6	-6	'2'		
'100007390148'	52	1	1	'3'	'2'	64101	670	13506	22515	36023	11165	1985	'1'	1	'-6'	-6	-6	'2'		
'100007540148'	46	3	3	'3'	'1'	61059	897	16501	27509	43993	17309	1985	'1'	1	'-6'	70000	-6	'1'		
'100008700141'	-9	2	-6	'4'	'4'	50820	743	11020	18370	29376	-6	1980	'3'	1	'-6'	-6	2	'-6'		
'100008960141'	56	3	2	'4'	'5'	51180	1129	12498	20829	33354	14353	1985	'1'	1	'-6'	195000	-6	'1'		

[illegible]

- ▼ Data Cleaning

1. Create new df containing only the chosen columns

```
df = raw_df[['AGE1', 'TENURE', 'FMTOWNRENT', 'BEDRMS', 'PER', 'ROOMS', 'NUNITS', 'OTHERCOST', 'ZINC2', 'BUILT', 'BURDEN', 'UTILITY', 'TOT$
df.show()
```

[illegible]

- ▼ Check for Duplicates

```
duplicated_count = df.groupby(df.columns).count().where(col("count") > 1).count()
```

```
print("There are {} duplicated rows in the dataset.".format(duplicated_count))
print("Number of rows before dropping duplicates is {}".format(df.count()))
```

```
# Drop duplicates and keep the first occurrence
df = df.dropDuplicates()
```

```
print("Number of rows after dropping duplicates is {}".format(df.count()))
```

```
There are 19 duplicated rows in the dataset.
Number of rows before dropping duplicates is 49090.
Number of rows after dropping duplicates is 49068.
```

```
df.describe().show()
```

summary	AGE1	TENURE	FMTOWNRENT	BEDRMS	PER	ROOMS	NUNITS	OT
count	49068	49068	49068	49068	49068	49068	49068	
mean	46.00014265916687	null	null	2.731515447949784	1.8160104344990626	5.659635607728051	10.844542267873155	65.1686299
stddev	23.310777795787647	null	null	1.074411489483994	2.7164447967521523	1.840251076342803	44.15603237738782	131.6596000
min	-9	'-6'	'1 Owner'	0	-6	1	-7	
max	93	'3'	'2 Renter'	10	14	21	981	5297

▼ Data Analysis

1. How many percentage are renter and owners in the whole dataset?


```
# Convert the Spark DataFrame to a Pandas DataFrame
pd_df1 = df.toPandas()

value_counts = pd_df1['FMTOWNRENT'].value_counts()

# Create a pie chart
plt.figure()
plt.pie(value_counts, labels=value_counts.index, explode=(0.1, 0.1), shadow=True, startangle=90, autopct=lambda p: '{:.2f}%\n({:.0f})'.format(p))

# Add a title to the chart
plt.title("Distribution of Renters and Owners")

# Show the pie chart
plt.show()
```



The majority of the dataset owned a house.

2. How many renters/owners based on their social class

An income quintile is a measure of neighbourhood socioeconomic status that divides the population into 5 income groups (from lowest income to highest income) so that approximately 20% of the population is in each group. So, we are taking the Household Income Distribution (HIQ) in the year 2009 as our dataset are from 1985-2009.

2006		2007		2008		2009		2010		
Quintile	Upper Limit	Mean	Upper Limit	Mean	Upper Limit	Mean	Upper Limit	Mean	Upper Limit	Mean
Lowest quintile	\$20,035	\$11,352	\$20,291	\$11,551	\$20,712	\$11,656	\$20,453	\$11,552	\$20,000	\$10,994
Second quintile	\$37,774	\$28,777	\$39,100	\$29,442	\$39,000	\$29,517	\$38,550	\$29,257	\$38,000	\$28,532
Middle quintile	\$60,000	\$48,223	\$62,000	\$49,968	\$62,725	\$50,132	\$61,801	\$49,534	\$61,500	\$49,167
Fourth quintile	\$97,032	\$76,329	\$100,000	\$79,111	\$100,240	\$79,760	\$100,000	\$78,694	\$100,029	\$78,877
Top quintile	--	\$168,170	--	\$167,971	--	\$171,057	--	\$170,844	--	\$169,391
Top 5% ¹	\$174,012	\$297,405	\$177,000	\$287,191	\$180,000	\$294,709	\$180,001	\$295,388	\$180,485	\$287,201

2001		2002		2003		2004		2005		
Quintile	Upper Limit	Mean	Upper Limit	Mean	Upper Limit	Mean	Upper Limit	Mean	Upper Limit	Mean
Lowest quintile	\$17,970	\$10,136	\$17,916	\$9,990	\$17,984	\$9,996	\$18,486	\$10,224	\$19,178	\$10,655
Second quintile	\$33,314	\$25,468	\$33,377	\$25,400	\$34,000	\$25,678	\$34,675	\$26,212	\$36,000	\$27,357
Middle quintile	\$53,000	\$42,629	\$53,162	\$42,802	\$54,453	\$43,588	\$55,230	\$44,411	\$57,660	\$46,301
Fourth quintile	\$83,500	\$66,839	\$84,016	\$67,326	\$86,867	\$68,994	\$88,002	\$70,026	\$91,705	\$72,825
Top quintile	--	\$145,970	--	\$143,743	--	\$147,078	--	\$151,438	--	\$159,583
Top 5% ¹	\$150,499	\$260,464	\$150,002	\$251,010	\$154,120	\$253,239	\$157,152	\$263,896	\$166,000	\$281,155

So, the division is as below

1. Lowest Quintile <20,034

2. 20,035 < Second Quintile < 38,549
3. 38,550 < Middle Quintile < 61,800
4. 61,801 < Fourth Quintile < 100,000
5. 100,001 < Top Quintile < 180,001
6. Top 5% > 180,001

Variables used

- Income = ZINC2
- Tenure = FMTOWNRENT

```
df2 = df[['ZINC2', 'FMTOWNRENT']]
df2.show()
```

```
+-----+-----+
| ZINC2|FMTOWNRENT|
+-----+-----+
| 70000|'2 Renter'|
| 20000|'1 Owner'|
| 6500|'2 Renter'|
|119000|'1 Owner'|
| 79900|'1 Owner'|
| 33000|'1 Owner'|
|108000|'1 Owner'|
| 16540|'1 Owner'|
| -6|'2 Renter'|
| 32400|'1 Owner'|
|101000|'1 Owner'|
| -6|'2 Renter'|
| 96900|'1 Owner'|
| 23800|'2 Renter'|
| 55000|'2 Renter'|
| -6|'2 Renter'|
|100000|'1 Owner'|
| 54900|'2 Renter'|
| 55000|'2 Renter'|
| 30000|'2 Renter'|
+-----+-----+
only showing top 20 rows
```

```
# Filter out rows with ZINC2 less than
df2 = df2.filter(df2['ZINC2'] > 0)
df2.describe().show()
```

```
+-----+-----+-----+
|summary|          ZINC2|FMTOWNRENT|
+-----+-----+-----+
| count|          44284|          44284|
| mean| 66996.3812663716|          null|
| stddev|68426.17762646428|          null|
| min|              1|'1 Owner'|
| max|          852840|'2 Renter'|
+-----+-----+-----+
```

```
from pyspark.sql import SparkSession
from pyspark.sql.functions import col, when
```

```
# Define the income ranges
income_ranges = [0, 20034, 38549, 61800, 100000, 180001, float('inf')]
income_range_labels = ['Lowest', 'Second', 'Middle', 'Fourth', 'Top 5%']
```

```
# Use the 'withColumn' function to add a new column with income ranges
df2 = df2.withColumn(
    'Quintile',
    when((col('ZINC2') >= income_ranges[0]) & (col('ZINC2') < income_ranges[1]), income_range_labels[0])
    .when((col('ZINC2') >= income_ranges[1]) & (col('ZINC2') < income_ranges[2]), income_range_labels[1])
    .when((col('ZINC2') >= income_ranges[2]) & (col('ZINC2') < income_ranges[3]), income_range_labels[2])
    .when((col('ZINC2') >= income_ranges[3]) & (col('ZINC2') < income_ranges[4]), income_range_labels[3])
    .when(col('ZINC2') >= income_ranges[4], income_range_labels[4])
    .otherwise('Unknown')
)
```

```
df2.show()
```

```
+-----+-----+-----+
| ZINC2|FMTOWNRENT|Quintile|
```

```

+-----+-----+-----+
| 70000|'2 Renter'| Fourth|
| 20000| '1 Owner'| Lowest|
| 6500| '2 Renter'| Lowest|
|119000| '1 Owner'| Top 5%|
| 79900| '1 Owner'| Fourth|
| 33000| '1 Owner'| Second|
|108000| '1 Owner'| Top 5%|
| 16540| '1 Owner'| Lowest|
| 32400| '1 Owner'| Second|
|101000| '1 Owner'| Top 5%|
| 96900| '1 Owner'| Fourth|
| 23800| '2 Renter'| Second|
| 55000| '2 Renter'| Middle|
|100000| '1 Owner'| Top 5%|
| 54900| '2 Renter'| Middle|
| 55000| '2 Renter'| Middle|
| 30000| '2 Renter'| Second|
|100000| '1 Owner'| Top 5%|
| 82000| '1 Owner'| Fourth|
| 74200| '1 Owner'| Fourth|
+-----+-----+-----+
only showing top 20 rows

```

```

# Drop a column
df2 = df2.drop('ZINC2')
df2.show()

```

```

+-----+-----+
|FMTOWNRENT|Quintile|
+-----+-----+
|'2 Renter'| Fourth|
| '1 Owner'| Lowest|
|'2 Renter'| Lowest|
| '1 Owner'| Top 5%|
| '1 Owner'| Fourth|
| '1 Owner'| Second|
| '1 Owner'| Top 5%|
| '1 Owner'| Lowest|
| '1 Owner'| Second|
| '1 Owner'| Top 5%|
| '1 Owner'| Fourth|
|'2 Renter'| Second|
|'2 Renter'| Middle|
| '1 Owner'| Top 5%|
|'2 Renter'| Middle|
|'2 Renter'| Middle|
|'2 Renter'| Second|
| '1 Owner'| Top 5%|
| '1 Owner'| Fourth|
| '1 Owner'| Fourth|
+-----+-----+
only showing top 20 rows

```

```

# Get unique values from the 'FMTOWNRENT' column
unique_fmtownrent_values = df2.select('FMTOWNRENT').distinct().rdd.flatMap(lambda x: x).collect()
print(unique_fmtownrent_values)

```

```

["'2 Renter'", "'1 Owner'"]

```

```

# Convert the Spark DataFrame to a Pandas DataFrame
pd_df2 = df2.toPandas()

```

```

# Filter the DataFrame to rows where 'FMTOWNRENT' is '1 Owner'
owner_df = pd_df2[pd_df2['FMTOWNRENT'] == "'1 Owner'"]

```

```

# Calculate the value counts for each unique value in the 'Quintile' column within '1 Owner' rows
owner_counts = owner_df['Quintile'].value_counts()

```

```

# Create a pie chart
plt.figure()
plt.pie(owner_counts, labels=owner_counts.index, shadow=True, startangle=90, autopct='%1.2f%%', explode=(0.1, 0.1, 0.1, 0.1, 0.1))

```

```


# Add a title to the chart
plt.title("Distribution of Owner by Quintile")

```

```

# Show the pie chart
plt.show()

```



About half of the owners are the fourth quintile and top 5%. This may be due to for rich people, property can be seen as an investment rather than their home.

```
# Convert the Spark DataFrame to a Pandas DataFrame
pd_df2 = df2.toPandas()


# Filter the DataFrame to rows where 'FMTOWNRENT' is '1 Owner'
renter_df = pd_df2[pd_df2['FMTOWNRENT'] == "'2 Renter'"]

# Calculate the value counts for each unique value in the 'Quintile' column within '1 Owner' rows
renter_counts = renter_df['Quintile'].value_counts()

# Create a pie chart
plt.figure()
plt.pie(renter_counts, labels=renter_counts.index, shadow=True, startangle=90, autopct='%1.2f%%', explode=(0.1, 0.1, 0.1, 0.1, 0.1))

# Add a title to the chart
plt.title("Distribution of Renter by Quintile")

# Show the pie chart
plt.show()
```



The Lowest & Second quintile dominates the rent's market by 2/3 of the data. This can be due to their low income makes it hard for them to buy a house (loan approval, house value, etc)

3. Do people with higher income have higher house value? (for both owner & renter)

Variables

- Income (ZINC2)
- House Value (VALUE)

```
df3 = df[['ZINC2', 'VALUE']]
df3.describe().show()
```

summary	ZINC2	VALUE
count	49068	49068
mean	60460.54775006114	158074.6332029021
stddev	67978.78827433457	250180.8188612107
min	-26976	-6
max	852840	2465647

```
# filter out income & house value less than zero
df3 = df3.filter(df3['ZINC2'] > 0)
df3 = df3.filter(df3['VALUE'] > 0)
```

```
df3.describe().show()
```

summary	ZINC2	VALUE
count	29905	29905
mean	79935.88426684501	246742.25326868417
stddev	75184.06437858166	273220.8036223053
min	1	1
max	852840	2465647

```
pd_df3 = df3.toPandas()
```


```
plt.figure(figsize=(10, 10))
```

```
# Scatter plot with enhanced visual features
plt.scatter(x=pd_df3['VALUE'], y=pd_df3['ZINC2'], color='blue', alpha=0.5, label='Data Points')
plt.title('Income vs. House Values')
plt.xlabel('House Value')
plt.ylabel('Income')
plt.grid(True, linestyle='--', alpha=0.7)
plt.legend()
```

```
# Limit axes range if needed
plt.xlim(0, 150000) # Adjust the values as necessary
plt.ylim(0, 900000) # Adjust the values as necessary
```

```
# Set custom x-axis tick locations and labels
custom_xticks = [0, 500000, 1000000, 1500000]
custom_xtick_labels = ['0', '0.5M', '1M', '1.5M']
plt.xticks(custom_xticks, custom_xtick_labels)
```

```
# Display the plot
plt.show()
```



```
plt.figure(figsize=(10, 10))
```

```
# Scatter plot with enhanced visual features
```

```
plt.scatter(x=pd_df3['VALUE'], y=pd_df3['ZINC2'], color='blue', alpha=0.5, label='Data Points')
```

```
plt.title('Income vs. House Values')
```

```
plt.xlabel('House Value')
```

```
plt.ylabel('Income')
```

```
plt.grid(True, linestyle='--', alpha=0.7)
```

```
plt.legend()
```

```
# Limit axes range if needed
```

```
plt.xlim(0, 1500000) # Adjust the values as necessary
```

```
plt.ylim(0, 500000) # Adjust the values as necessary
```

```
# Set custom x-axis tick locations and labels
```


```
custom_xticks = [0, 500000, 1000000, 1500000]
```

```
custom_xtick_labels = ['0', '0.5M', '1M', '1.5M']
```

```
plt.xticks(custom_xticks, custom_xtick_labels)
```

```
# Display the plot
```

```
plt.show()
```

4. Do number of people in a household increase as the age of the head of a household increase?

Variables

- No of people in a household = PER
- Age of the head of household = AGE1

```
df4 = df[['PER', 'AGE1']]
df4.describe().show()
```

summary	PER	AGE1
count	49068	49068
mean	1.8160104344990626	46.00014265916687
stddev	2.7164447967521523	23.310777795787647
min	-6	-9
max	14	93

```
# Filter negative data
```

```
df4 = df4.filter(df4['PER'] > 0)
df4 = df4.filter(df4['AGE1'] > 0)
```

```
df4.describe().show()
```

summary	PER	AGE1
count	45057	45057
mean	2.5117961692966686	50.896286925449985
stddev	1.4537738374258677	17.277128359384527
min	1	14
max	14	93

```
pd_df4 = df4.toPandas()
```

```
# Create age bins (you can adjust the bin size as needed)
```

```
bin_size = 10
```

```
age_bins = np.arange(0, pd_df4['AGE1'].max() + bin_size, bin_size)
```

```
# Cut 'AGE1' values into bins
```

```
pd_df4['AgeBin'] = pd.cut(pd_df4['AGE1'], bins=age_bins, right=False)
```

```
# Group by 'AgeBin' and calculate the mean 'PER' for each bin
```

```
age_per_mean = pd_df4.groupby('AgeBin')['PER'].mean()
```

```
# Create a bar graph
```

```
plt.bar(age_per_mean.index.astype(str), age_per_mean)
```

```
plt.title('Average PER by Age Bin')
```

```
plt.xlabel('Age')
```

```
plt.ylabel('Average PER')
```

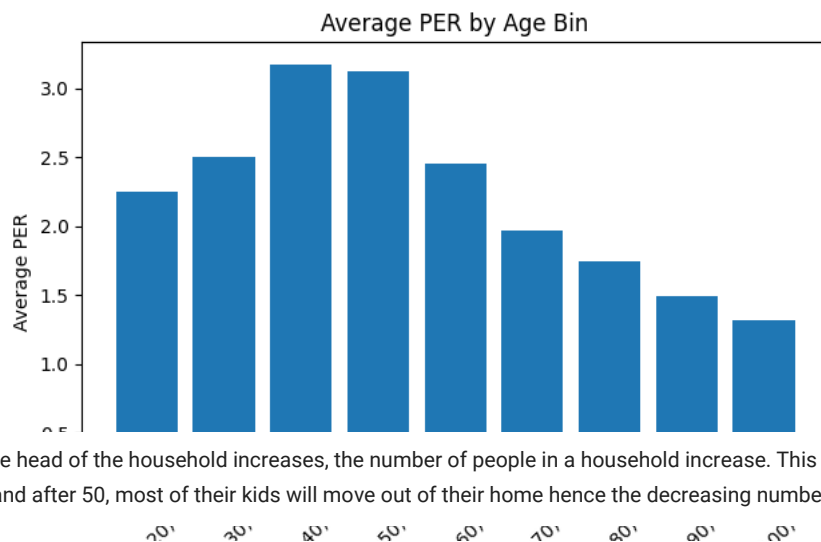
```
plt.xticks(rotation=45)
```

```
plt.tight_layout()
```

```
# Show the bar graph
```

```
plt.show()
```





As the head of the household increases, the number of people in a household increase. This can be largely due to they having kids during that age and after 50, most of their kids will move out of their home hence the decreasing number of Average people in a household.

5. Do older people prefer single unit or building with multiple units?

Variables

- Units = STRUCTURETYPE
- Older people = AGE1

```
df5 = df[['AGE1', 'STRUCTURETYPE']]
df5.describe().show()
```

summary	AGE1	STRUCTURETYPE
count	49068	49068
mean	46.00014265916687	1.819821472242602
stddev	23.310777795787647	1.4424372779554488
min	-9	-9
max	93	6

```
df5 = df5.filter(df5['AGE1'] > 60)
df5 = df5.withColumn('STRUCTURETYPE', when(col('STRUCTURETYPE') == 1, col('STRUCTURETYPE')).otherwise(2))
```

```
df5.describe().show()
```

summary	AGE1	STRUCTURETYPE
count	12974	12974
mean	72.54208416833667	1.2619855094804995
stddev	8.653217421555254	0.4397317435266575
min	61	1
max	93	2

```
pd_df5 = df5.toPandas()
# Define age ranges
age_bins = np.arange(60, 101, 10)
age_labels = [f'{start}-{end-1}' for start, end in zip(age_bins, age_bins + 10)]

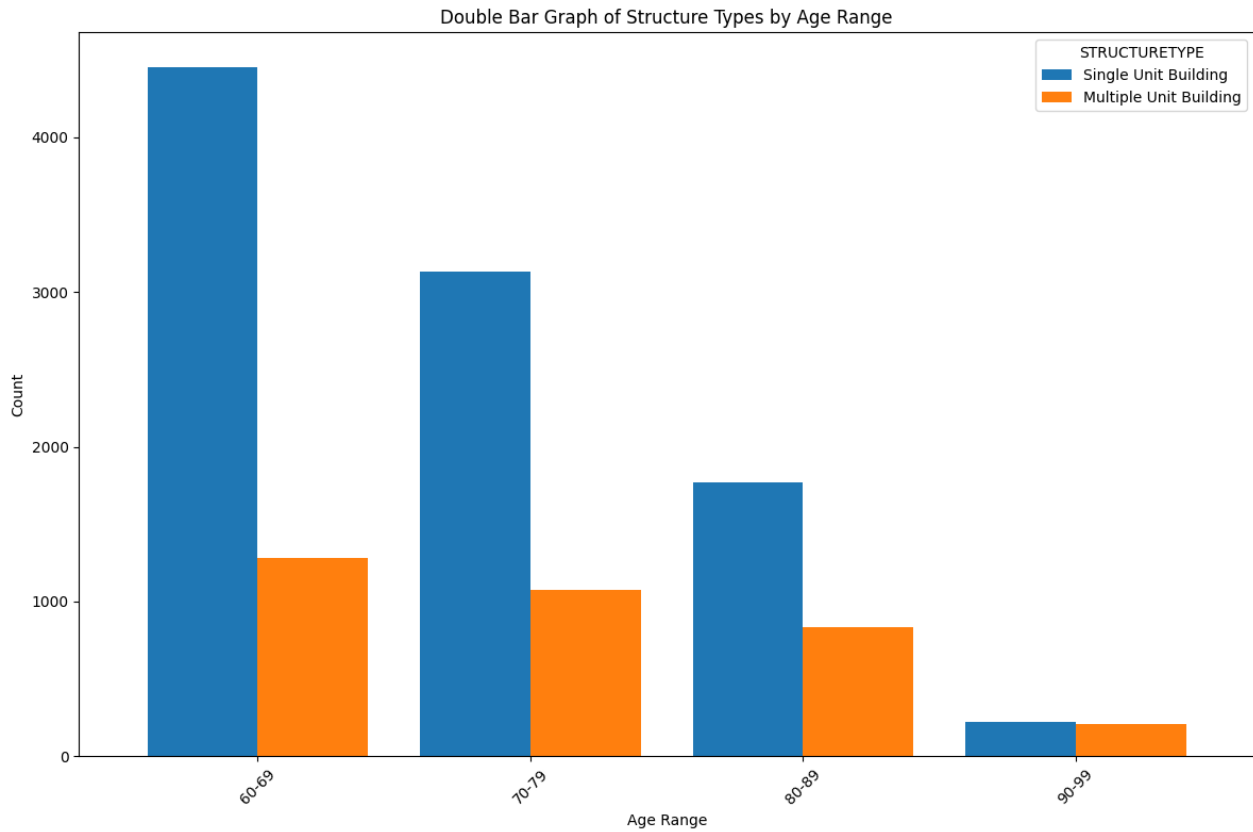
# Cut 'AGE1' values into age ranges
pd_df5['AgeRange'] = pd.cut(pd_df5['AGE1'], bins=age_bins, labels=age_labels[:-1], right=False)

# Group by 'AgeRange' and 'STRUCTURETYPE', and calculate value counts
age_structure_counts = pd_df5.groupby(['AgeRange', 'STRUCTURETYPE']).size().unstack()

# Create a double bar graph
ax = age_structure_counts.plot(kind='bar', stacked=False, figsize=(12, 8), width=0.81)
plt.xlabel('Age Range')
plt.ylabel('Count')
plt.title('Double Bar Graph of Structure Types by Age Range')
plt.legend(title='STRUCTURETYPE', labels=['Single Unit Building', 'Multiple Unit Building'])

plt.xticks(rotation=45)
plt.tight_layout()
```

```
plt.show()
```



For a while, older people until the age of 80 prefer to live alone. But when they are 90 and above, they tend to live in a community. This can be due to their family encourage them not to live alone because it is hard for them to know if anything happen to their parents

6. Do ZADEQ affect the value of the house and Fair Market Rent?

According to [American Housing Survey](#), The Adequacy Index (named ZADEQ prior to the 2015 AHS) is comprised of eight criterion capturing severe physical deficiencies like no running water, plumbing, etc. US Department of Housing and Urban Development (HUD) created this measure to assess the extent to which the housing stock met the standard of “a decent home and a suitable living environment,”

The HUD measure is contained in the variable ZADEQ, which takes three values:

- 1 if a unit is considered adequate,
- 2 if a unit is considered moderately inadequate, and
- 3 if a unit is considered severely inadequate

Variables

- Adequacy Index = ZADEQ
- Value = VALUE, FMR

```
df6 = df[['ZADEQ', 'VALUE', 'FMR']]
df6.describe().show()
```

	summary	ZADEQ	VALUE	FMR
count	49068	49068	49068	49068
mean	null	158074.6332029021	1073.141171435588	
stddev	null	250180.8188612107	355.65992591671267	
min	'-6'	-6	427	
max	'3'	2465647	3501	

```

print("Number of rows before filtering:", df6.count())
# Filter out rows with ZADEQ values less than 0
df6 = df6.filter(df6['ZADEQ'] != "-6")
print("Number of rows after filtering:", df6.count())
# Show the resulting DataFrame
df6.describe().show()

```

```

Number of rows before filtering: 49068
Number of rows after filtering: 45057
+-----+-----+-----+-----+
|summary|ZADEQ|          VALUE|          FMR|
+-----+-----+-----+-----+
| count|45057|          45057|          45057|
|  mean| null| 165356.5500366203|1083.2290654060412|
| stddev| null|252187.55376145715|357.53157953189094|
|   min|  '1'|             -6|             427|
|   max|  '3'|          2465647|            3501|
+-----+-----+-----+-----+

```

```

plt.figure(figsize=(10, 10))
pd_df6 = df6.toPandas()


# Scatter plot with enhanced visual features
plt.scatter(x=pd_df6['ZADEQ'], y=pd_df6['VALUE'], color='blue', alpha=0.5, label='Data Points')
plt.title('Adequacy Index vs. House Values')
plt.xlabel('ZADEQ')
plt.ylabel('House Value')
plt.grid(True, linestyle='--', alpha=0.7)
plt.legend()

# Limit axes range if needed
plt.ylim(0, 500000) # Adjust the values as necessary

# Set custom x-axis tick locations and labels
custom_yticks = [0, 500000, 1000000, 1500000]
custom_ytick_labels = ['0', '0.5M', '1M', '1.5M']
plt.yticks(custom_yticks, custom_ytick_labels)

# Display the plot
plt.show()

```




```
plt.figure(figsize=(10, 10))
pd_df6 = df6.toPandas()

# Scatter plot with enhanced visual features
plt.scatter(x=pd_df6['ZADEQ'], y=pd_df6['FMR'], color='blue', alpha=0.5, label='Data Points')
plt.title('Adequacy Index vs. Fair Market Rent')
plt.xlabel('ZADEQ')
plt.ylabel('Fair Market Rent')
plt.grid(True, linestyle='--', alpha=0.7)
plt.legend()

# Limit axes range if needed
#plt.ylim(0, 500000) # Adjust the values as necessary

# Set custom x-axis tick locations and labels
#custom_yticks = [0, 500000, 1000000, 1500000]
#custom_ytick_labels = ['0', '0.5M', '1M', '1.5M']
#plt.yticks(custom_yticks, custom_ytick_labels)

# Display the plot
plt.show()
```



Both the value and Fair Market Rent are higher if the score fo Adequacy Index is high. This can be due to people not wanting to live in an inadequate house wiich drive the value down due to poor demand/

7. Do median income in an area affect what type of home structure that they built?

Variables

- Area median income = LMED
- Home structure = STRUCTURETYPE

```
df7 = df[['LMED', 'STRUCTURETYPE']]
df7.describe().show()
```

```
+-----+-----+-----+
|summary|          LMED|    STRUCTURETYPE|
+-----+-----+-----+
| count|          49068|          49068|
| mean| 66716.96500774435| 1.819821472242602|
| stddev|11943.529169776679| 1.4424372779554488|
| min|          32000|          -9|
| max|         122300|           6|
+-----+-----+-----+
```

```
df7 = df7.filter(df7['STRUCTURETYPE'] > 0)
df7.describe().show()
```

```
+-----+-----+-----+
|summary|          LMED|    STRUCTURETYPE|
+-----+-----+-----+
| count|          49060|          49060|
| mean| 66717.6581736649| 1.8215858132898493|
| stddev|11944.016016223557| 1.435921734985718|
| min|          32000|           1|
| max|         122300|           6|
+-----+-----+-----+
```

```
# Define the income ranges
income_ranges = [0, 20034, 38549, 61800, 100000, 180001, float('inf')]
income_range_labels = ['Lowest', 'Second', 'Middle', 'Fourth', 'Top 5%']
```

```
# Use the 'withColumn' function to add a new column with income ranges
df7 = df7.withColumn('Quintile',
  when((col('LMED') >= income_ranges[0]) & (col('LMED') < income_ranges[1]), income_range_labels[0])
  .when((col('LMED') >= income_ranges[1]) & (col('LMED') < income_ranges[2]), income_range_labels[1])
  .when((col('LMED') >= income_ranges[2]) & (col('LMED') < income_ranges[3]), income_range_labels[2])
  .when((col('LMED') >= income_ranges[3]) & (col('LMED') < income_ranges[4]), income_range_labels[3])
  .when(col('LMED') >= income_ranges[4], income_range_labels[4])
  .otherwise('Unknown')
)
```

```
df7.describe().show()
```

```
+-----+-----+-----+-----+
|summary|          LMED|    STRUCTURETYPE|Quintile|
+-----+-----+-----+-----+
| count|          49060|          49060| 49060|
| mean| 66717.6581736649| 1.8215858132898493| null|
| stddev|11944.016016223557| 1.435921734985718| null|
| min|          32000|           1| Fourth|
| max|         122300|           6| Top 5%|
+-----+-----+-----+-----+
```

```
# Find the unique values in the Quintile column
unique_quintiles = df7.select('Quintile').distinct()
unique_quintiles.show()
pd_df7 = df7.toPandas()
```

```
+-----+
|Quintile|
+-----+
| Fourth|
| Middle|
| Second|
| Top 5%|
+-----+
```

```

# Filter the DataFrame for the Second "Quintile"
second_quintile_df = pd_df7[pd_df7['Quintile'] == 'Second']

# Define the structure type labels
structure_type_labels = {
    1: 'Single unit',
    2: '2-4 units',
    3: '5-19 units',
    4: '20-49 units',
    5: '50+ units',
    6: 'Mobile homes'
}

# Count the occurrences of each STRUCTURETYPE within the Second "Quintile"
structure_counts = second_quintile_df['STRUCTURETYPE'].value_counts()

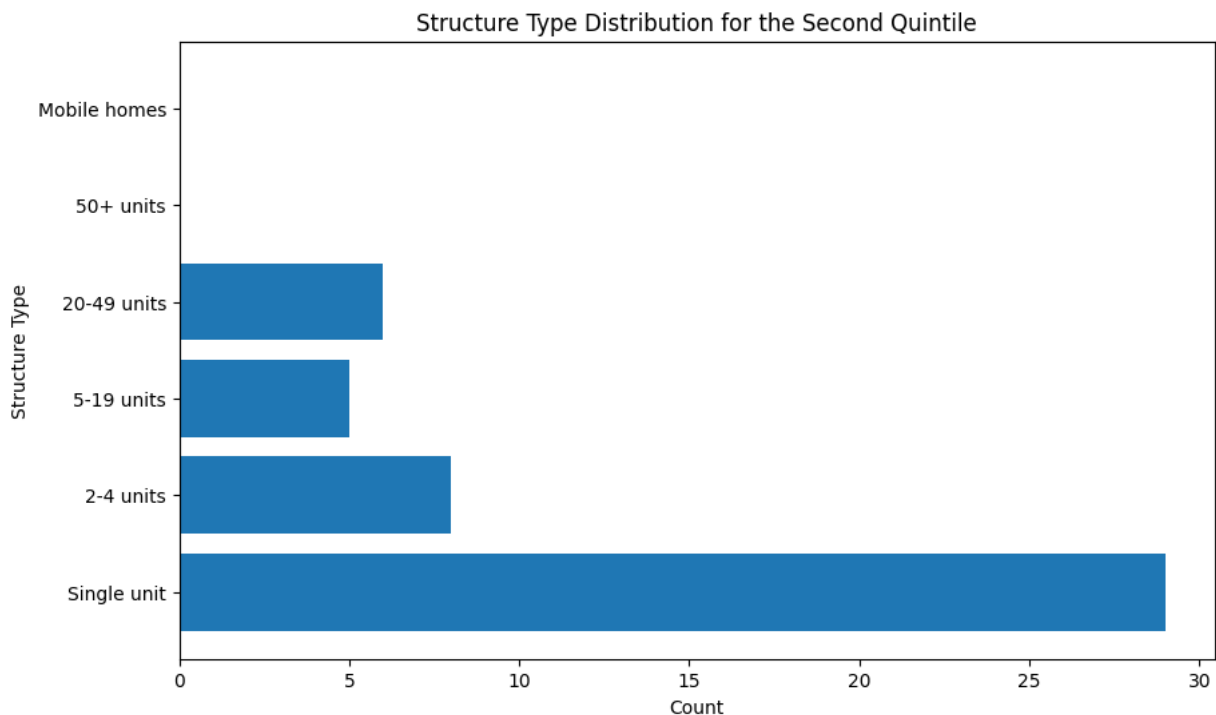
# Create a Series with all possible STRUCTURETYPE values and fill with zeros
all_structure_types = pd.Series(index=range(1, 7), data=0)
all_structure_types.update(structure_counts) # Update with actual counts

# Create the bar plot
plt.figure(figsize=(10, 6))

# Plotting all possible values 1 to 6 with their corresponding counts
plt.barh(all_structure_types.index, all_structure_types.values, tick_label=[structure_type_labels[i] for i in all_structure_types.index])
plt.xlabel('Count')
plt.ylabel('Structure Type')
plt.title('Structure Type Distribution for the Second Quintile')

plt.show()

```




```

# Filter the DataFrame for the Fourth "Quintile"
fourth_quintile_df = pd_df7[pd_df7['Quintile'] == 'Fourth']

# Count the occurrences of each STRUCTURETYPE within the Fourth "Quintile"
structure_counts = fourth_quintile_df['STRUCTURETYPE'].value_counts()

# Create the bar plot
plt.figure(figsize=(10, 6))
plt.barh(structure_counts.index, structure_counts.values, tick_label=[structure_type_labels[i] for i in all_structure_types.index])
plt.xlabel('Count')
plt.ylabel('Structure Type')
plt.title('Structure Type Distribution for the Fourth Quintile')
plt.show()


```



```
# Filter the DataFrame for the Fourth "Quintile"
top5_quintile_df = pd_df7[pd_df7['Quintile'] == 'Top 5%']

# Count the occurrences of each STRUCTURETYPE within the Top 5% "Quintile"
structure_counts = top5_quintile_df['STRUCTURETYPE'].value_counts()

# Create the bar plot
plt.figure(figsize=(10, 6))
plt.barh(structure_counts.index, structure_counts.values, tick_label=[structure_type_labels[i] for i in all_structure_types.index])
plt.xlabel('Count')
plt.ylabel('Structure Type')
plt.title('Structure Type Distribution for the Top 5%')
plt.show()
```



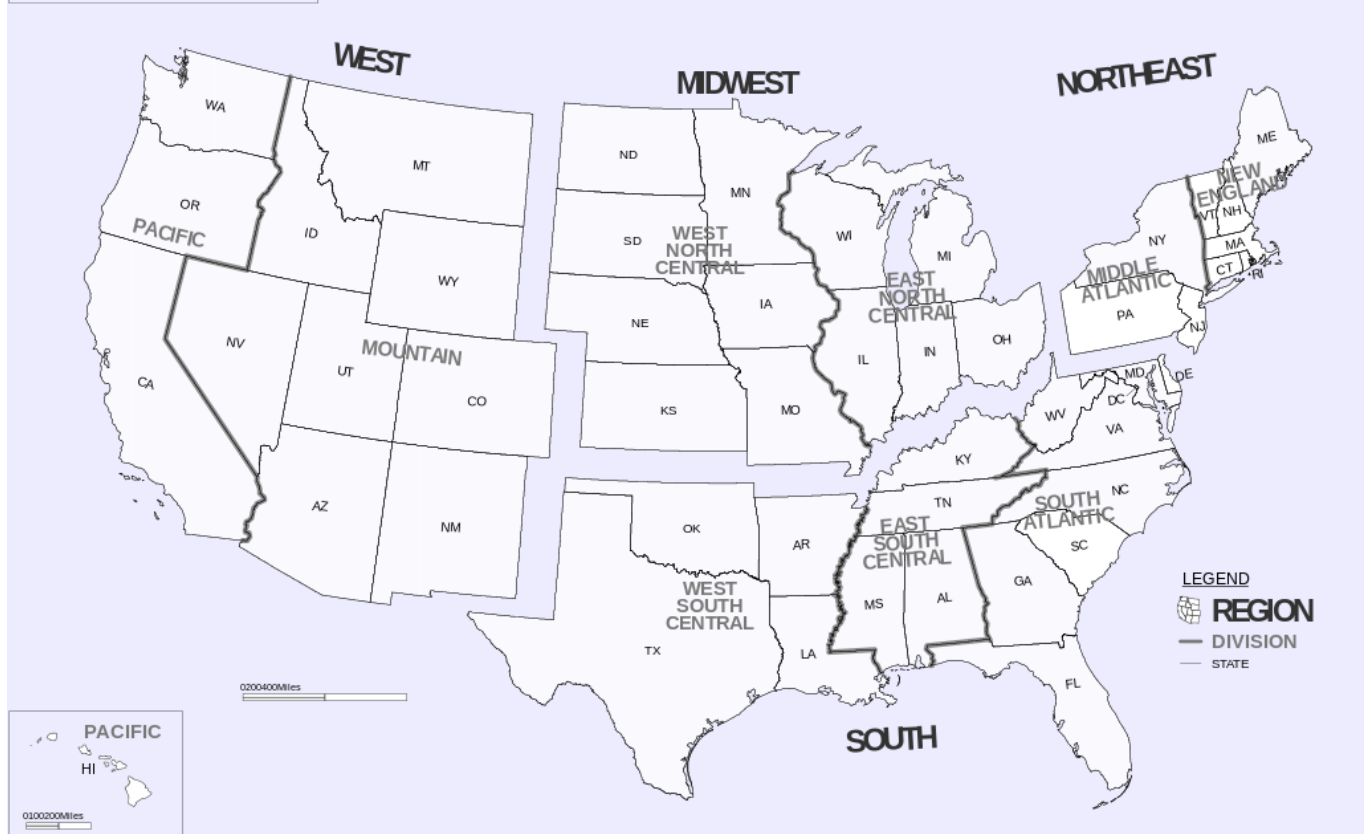
Analysis

8. In each region, what is the most popular type of structure?

Since 1950, the United States Census Bureau defines four statistical regions, with nine divisions. The Census Bureau region definition is "widely used for data collection and analysis", and is the most commonly used classification system. [Wikipedia](https://www.census.gov/programs-surveys/geography/about.html)



Census Regions and Divisions of the United States



- Region 1 - **Northeast** (massachusetts, rhode island, etc)
- Region 2 - **Midwest** (michigan, ohio, nebraska)
- Region 3 - **South** (Washington DC, texas)
- Region 4 - **West** (california, arizona)

Variables

- Location = REGION
- Home structure = STRUCTURETYPE

```
df8 = df['REGION', 'STRUCTURETYPE']
df8.describe().show()
```

```
+-----+-----+
|summary|REGION|  STRUCTURETYPE|
+-----+-----+
| count| 49068|          49068|
| mean| null| 1.819821472242602|
| stddev| null| 1.4424372779554488|
| min| '1'|          -9|
| max| '4'|          6|
+-----+-----+
```

```
df8 = df8.filter(df7['STRUCTURETYPE']>0)
df8.describe().show()
pd_df8 = df8.toPandas()
```

```
+-----+-----+
|summary|REGION|  STRUCTURETYPE|
+-----+-----+
| count| 49060|          49060|
| mean| null| 1.8215858132898493|
| stddev| null| 1.435921734985718|
| min| '1'|          1|
| max| '4'|          6|
+-----+-----+
```

```

# Define the structure type labels and colors
structure_type_labels = {
    1: 'Single unit',
    2: '2-4 units',
    3: '5-19 units',
    4: '20-49 units',
    5: '50+ units',
    6: 'Mobile homes'
}

structure_type_colors = {
    1: 'blue',
    2: 'orange',
    3: 'green',
    4: 'red',
    5: 'purple',
    6: 'brown'
}

# Group the DataFrame by REGION and STRUCTURETYPE, and count occurrences
grouped = pd_df8.groupby(['REGION', 'STRUCTURETYPE']).size().reset_index(name='COUNT')

# Get unique regions and structure types
unique_regions = grouped['REGION'].unique()
unique_structure_types = grouped['STRUCTURETYPE'].unique()

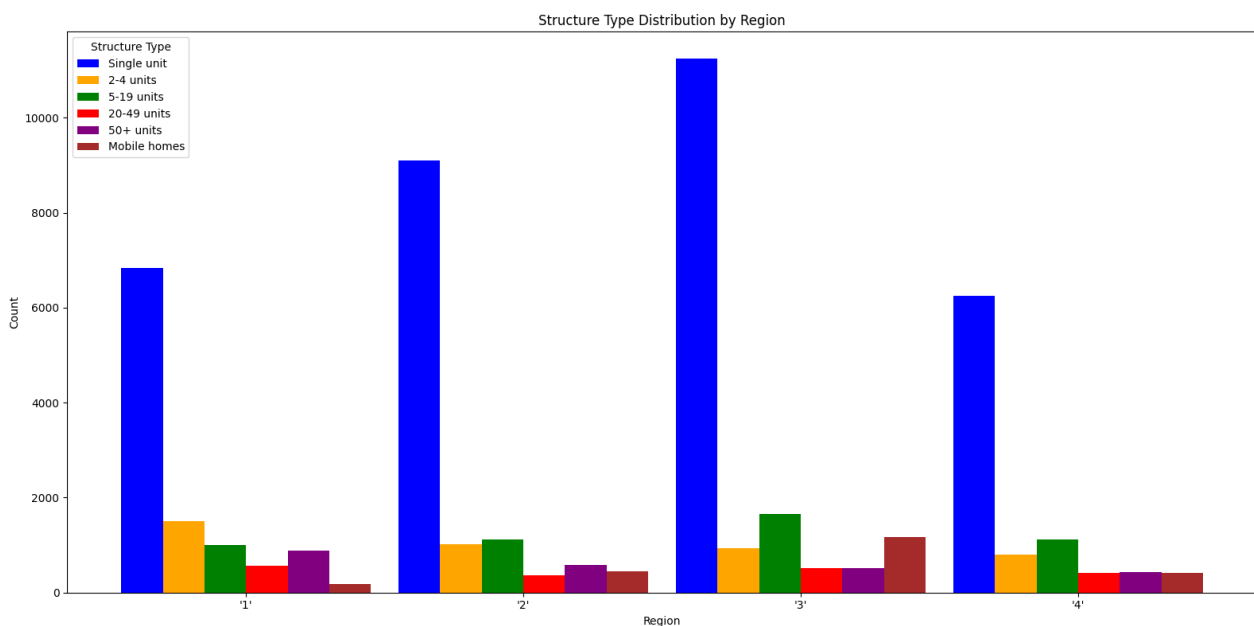
# Create the main plot
plt.figure(figsize=(16, 8))

# Iterate through each structure type and plot grouped bars for each region
for i, st in enumerate(unique_structure_types):
    structure_data = grouped[grouped['STRUCTURETYPE'] == st]
    counts = structure_data['COUNT']
    x = np.arange(len(unique_regions)) + i * 0.15
    plt.bar(x, counts, width=0.15, color=structure_type_colors[st], label=structure_type_labels[st])

# Set labels, title, and legend
plt.xlabel('Region')
plt.ylabel('Count')
plt.title('Structure Type Distribution by Region')
plt.xticks(np.arange(len(unique_regions)) + (len(unique_structure_types) - 1) * 0.15 / 2, unique_regions)
plt.legend(title='Structure Type', loc='upper left')

plt.tight_layout()
plt.show()

```



The majority of homes are single unit which is not surprising considering the US is wide and big so they would much prefer single unit homes.

In this context, the high number of multi-unit buildings in Region 1 aligns with the concentration of densely populated areas. Metropolises like New York City within this region are famously known for their towering apartment buildings that utilize available space to house a significant number of residents.

9. Do region correspond to the market value of the unit, based on the number of rooms in the unit?

Variables

- Location = REGION
- Market Value of a unit = VALUE
- Number of rooms in a unit = NUNITS

```
df9 = df[['REGION', 'VALUE', 'NUNITS']]
df9.describe().show()
```

	summary	REGION	VALUE	NUNITS
count	49068		49068	49068
mean	null	158074.6332029021	10.844542267873155	
stddev	null	250180.8188612107	44.15603237738782	
min	'1'		-6	-7
max	'4'		2465647	981

```
df9 = df9.filter(df9['VALUE'] > 0)
df9 = df9.filter(df9['NUNITS'] > 0)
```

```
df9.describe().show()
```

	summary	REGION	VALUE	NUNITS
count	31477		31477	31477
mean	null	246418.38955427773	3.4509641960796773	
stddev	null	275317.3265830988	24.718178699429878	
min	'1'		1	1
max	'4'		2465647	981

```
pd_df9 = df9.toPandas()
```

```
# Define colors for each region
```

```
region_colors = {
    "'1'": 'blue',
    "'2'": 'orange',
    "'3'": 'green',
    "'4'": 'red'
}
```

```
# Create the scatter plot
```

```
plt.figure(figsize=(10, 6))
```

```
for region, color in region_colors.items():
    region_subset = pd_df9[pd_df9['REGION'] == region]
    plt.scatter(region_subset['NUNITS'], region_subset['VALUE'], color=color, label=f'Region {region}')
```

```
# Set labels, title, and legend
```


```
plt.xlabel('Number of Units in Building')
```

```
plt.ylabel('Market Value of Unit')
```

```
plt.title('Relationship between Number of Units and Market Value by Region')
```

```
plt.legend(title='Region')
```

```
plt.show()
```



10. Are people segregated by the status of their income?

Do rich people only live with other rich people and poor people only lives with poor people?

Variables

- Income = Quintile & ZINC2
- Median income of an area = LMED

```
df10 = df[['LMED', 'ZINC2', 'AGE1']]
```

```
df10 = df10.filter(df10['LMED'] > 0)
df10 = df10.filter(df10['ZINC2'] > 0)
```

```
df10.describe().show()
```

summary	LMED	ZINC2	AGE1
count	44284	44284	44284
mean	66768.09703278837	66996.3812663716	51.030756029265646
stddev	11998.300086330291	68426.17762646428	17.23329424953238
min	32000	1	14
max	122300	852840	93

```
from pyspark.sql import SparkSession
from pyspark.sql.functions import col, when
```

```
# Define the income ranges
income_ranges = [0, 20034, 38549, 61800, 100000, 180001, float('inf')]
income_range_labels = ['Lowest', 'Second', 'Middle', 'Fourth', 'Top 5%']
```

```
# Use the 'withColumn' function to add a new column with income ranges
df10 = df10.withColumn(
    'Quintile',
    when((col('ZINC2') >= income_ranges[0]) & (col('ZINC2') < income_ranges[1]), income_range_labels[0])
    .when((col('ZINC2') >= income_ranges[1]) & (col('ZINC2') < income_ranges[2]), income_range_labels[1])
    .when((col('ZINC2') >= income_ranges[2]) & (col('ZINC2') < income_ranges[3]), income_range_labels[2])
    .when((col('ZINC2') >= income_ranges[3]) & (col('ZINC2') < income_ranges[4]), income_range_labels[3])
    .when(col('ZINC2') >= income_ranges[4], income_range_labels[4])
    .otherwise('Unknown')
)
```

```
df10.describe().show()
```

summary	LMED	ZINC2	AGE1	Quintile
count	44284	44284	44284	44284
mean	66768.09703278837	66996.3812663716	51.030756029265646	null
stddev	11998.300086330291	68426.17762646428	17.23329424953238	null
min	32000	1	14	Fourth
max	122300	852840	93	Top 5%

```
pd_df10 = df10.toPandas()
```

```
# Find the unique values in the Quintile column
unique_quintiles = df10.select('Quintile').distinct()
unique_quintiles.show()
```

```
+-----+
|Quintile|
+-----+
|  Fourth|
|  Middle|
|  Second|
|  Lowest|
|  Top 5%|
+-----+
```

```
# Filter the DataFrame for 'Lowest' quintile
lowest_quintile_df = pd_df10[pd_df10['Quintile'] == 'Lowest']
```

```
# Calculate the average LMED
average_income = lowest_quintile_df['ZINC2'].mean()
```

```
# Create the scatter plot
plt.figure(figsize=(10, 6))
```


```
# Scatter plot for 'Lowest' quintile
plt.scatter(lowest_quintile_df['AGE1'], lowest_quintile_df['ZINC2'], label='Lowest Quintile')
```

```
# Line for average LMED
plt.axhline(y=average_income, color='red', linestyle='--', label='Average Income')
```

```
# Set labels, title, and legend
plt.xlabel('Age')
plt.ylabel('Income')
plt.title('Scatter Plot of Income vs Age for Lowest Quintile')
plt.legend()
```

```
# Set y-axis limits
plt.ylim(0, 25000)
```

```
plt.show()
```



Scatter Plot of Income vs. Age for the Lowest Quintile appears relatively uniform. This uniformity could be due to the fact that individuals within the lowest quintile tend to have more consistent income and age demographics, resulting in less variation in the data points.

```
top5_quintile_df = pd_df10[pd_df10['Quintile'] == 'Top 5%']
```

```

# Calculate the average LMED
average_income = top5_quintile_df['ZINC2'].mean()

# Create the scatter plot
plt.figure(figsize=(10, 6))

# Scatter plot for 'Lowest' quintile
plt.scatter(top5_quintile_df['AGE1'], top5_quintile_df['ZINC2'], label='Top 5% Quintile')


# Line for average LMED
plt.axhline(y=average_income, color='red', linestyle='--', label='Average Income')

# Set labels, title, and legend
plt.xlabel('Age')
plt.ylabel('Income')
plt.title('Scatter Plot of Income vs Age for Top 5% Quintile')
plt.legend()

# Set y-axis limits
#plt.ylim(0, 250000)

plt.show()

```



This shows that for the top 5%, the income ranges are so wide that 800k income-person can live with 100k income-person.

A wide income range within the top 5% quintile, which can indeed result in individuals with significantly different income levels falling within the same category. This variation could be due to various factors, such as differences in household composition, location, or other socio-economic factors

